

```
In [64]: import pandas as pd
import numpy as np

import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score

from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
!pip install xgboost
from xgboost import XGBClassifier


from sklearn.preprocessing import StandardScaler
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.impute import SimpleImputer

!pip install nbconvert[webpdf]
!playwright install
```

```
Requirement already satisfied: xgboost in c:\users\pablo\appdata\local\abf certificate\python\python\lib\site-packages (3.1.2)
Requirement already satisfied: numpy in c:\users\pablo\appdata\local\abf certificate\python\python\lib\site-packages (from xgboost) (1.26.1)
Requirement already satisfied: scipy in c:\users\pablo\appdata\local\abf certificate\python\python\lib\site-packages (from xgboost) (1.11.3)
Requirement already satisfied: nbconvert[webpdf] in c:\users\pablo\appdata\local\abf certificate\python\python\lib\site-packages (7.16.6)
Requirement already satisfied: beautifulsoup4 in c:\users\pablo\appdata\local\abf certificate\python\python\lib\site-packages (from nbconvert[webpdf]) (4.14.2)
Requirement already satisfied: bleach!=5.0.0 in c:\users\pablo\appdata\local\abf certificate\python\python\lib\site-packages (from bleach[css]!=5.0.0->nbconvert[webpdf]) (6.2.0)
Requirement already satisfied: defusedxml in c:\users\pablo\appdata\local\abf certificate\python\python\lib\site-packages (from nbconvert[webpdf]) (0.7.1)
Requirement already satisfied: jinja2>=3.0 in c:\users\pablo\appdata\local\abf certificate\python\python\lib\site-packages (from nbconvert[webpdf]) (3.1.6)
Requirement already satisfied: jupyter-core>=4.7 in c:\users\pablo\appdata\local\abf certificate\python\python\lib\site-packages (from nbconvert[webpdf]) (5.9.1)
Requirement already satisfied: jupyterlab-pygments in c:\users\pablo\appdata\local\abf certificate\python\python\lib\site-packages (from nbconvert[webpdf]) (0.3.0)
Requirement already satisfied: markupsafe>=2.0 in c:\users\pablo\appdata\local\abf certificate\python\python\lib\site-packages (from nbconvert[webpdf]) (3.0.3)
Requirement already satisfied: mistune<4,>=2.0.3 in c:\users\pablo\appdata\local\abf certificate\python\python\lib\site-packages (from nbconvert[webpdf]) (3.1.4)
Requirement already satisfied: nbclient>=0.5.0 in c:\users\pablo\appdata\local\abf certificate\python\python\lib\site-packages (from nbconvert[webpdf]) (0.10.2)
Requirement already satisfied: nbformat>=5.7 in c:\users\pablo\appdata\local\abf certificate\python\python\lib\site-packages (from nbconvert[webpdf]) (5.10.4)
Requirement already satisfied: packaging in c:\users\pablo\appdata\local\abf certificate\python\python\lib\site-packages (from nbconvert[webpdf]) (25.0)
Requirement already satisfied: pandocfilters>=1.4.1 in c:\users\pablo\appdata\local\abf certificate\python\python\lib\site-packages (from nbconvert[webpdf]) (1.5.1)
Requirement already satisfied: pygments>=2.4.1 in c:\users\pablo\appdata\local\abf certificate\python\python\lib\site-packages (from nbconvert[webpdf]) (2.19.2)
Requirement already satisfied: traitlets>=5.1 in c:\users\pablo\appdata\local\abf certificate\python\python\lib\site-packages (from nbconvert[webpdf]) (5.14.3)
Collecting playwright (from nbconvert[webpdf])
  Downloading playwright-1.57.0-py3-none-win_amd64.whl.metadata (3.5 kB)
Requirement already satisfied: webencodings in c:\users\pablo\appdata\local\abf certificate\python\python\lib\site-packages (from bleach!=5.0.0->bleach[css]!=5.0.0->nbconvert[webpdf]) (0.5.1)
Requirement already satisfied: tinycc2<1.5,>=1.1.0 in c:\users\pablo\appdata\local\abf certificate\python\python\lib\site-packages (from bleach[css]!=5.0.0->nbconvert[webpdf]) (1.4.0)
Requirement already satisfied: platformdirs>=2.5 in c:\users\pablo\appdata\local\abf certificate\python\python\lib\site-packages (from jupyter-core>=4.7->nbconvert[webpdf]) (4.5.0)
Requirement already satisfied: jupyter-client>=6.1.12 in c:\users\pablo\appdata\local\abf certificate\python\python\lib\site-packages (from nbclient>=0.5.0->nbconvert[webpdf]) (8.6.3)
Requirement already satisfied: python-dateutil>=2.8.2 in c:\users\pablo\appdata\local\abf certificate\python\python\lib\site-packages (from jupyter-client>=6.1.12->nbclient>=0.5.0->nbconvert[webpdf]) (2.9.0.post0)
Requirement already satisfied: pyzmq>=23.0 in c:\users\pablo\appdata\local\abf certificate\python\python\lib\site-packages (from jupyter-client>=6.1.12->nbclient>=0.5.0
```



```
Successfully installed greenlet-3.3.0 playwright-1.57.0 pyee-13.0.0
Downloading Chromium 143.0.7499.4 (playwright build v1200) from https://cdn.playwright.dev/dbazure/download/playwright/builds/chromium/1200/chromium-win64.zip
```

|

0% of 169.8 MiB
|â- â- â- â- â- â-
| 10% of 169.8 MiB
|â- â-
| 20% of 169.8 MiB
|â- â-
| 30% of 169.8 MiB
|â- â-
| 40% of 169.8 MiB
|â- â-
| 50%
of 169.8 MiB
|â- â- â- â- â- â- â- â- â- â-
| 60% of 169.8 MiB
|â- â- â- â- â- â- â- â- â-
| 70% of 169.8 MiB
|â- â- â- â- â- â- â- â-
| 80% of 169.8 MiB
|â- â- â- â- â- â- â- â-
| 90% of 169.8 MiB
|â- â- â- â- â- â- â- â-
| 100% of 169.8 MiB
69.8 MiB

Chromium 143.0.7499.4 (playwright build v1200) downloaded to C:\Users\pablo\AppData\Local\ms-playwright\chromium-1200

Downloading Chromium Headless Shell 143.0.7499.4 (playwright build v1200) from http://cdn.playwright.dev/dbazure/download/playwright/builds/chromium/1200/chromium-headless-shell-win64.zip

|

0% of 107.2 MiB
|â- â- â- â- â- â-
| 10% of 107.2 MiB
|â- â- â- â- â- â- â-
| 20% of 107.2 MiB
|â- â- â- â- â- â- â-
| 30% of 107.2 MiB
|â- â- â- â- â- â- â-
| 40% of 107.2 MiB
|â- â- â- â- â- â- â-
| 50%
of 107.2 MiB
|â- â- â- â- â- â- â-
| 60% of 107.2 MiB
|â- â- â- â- â- â- â-
| 70% of 107.2 MiB
|â- â- â- â- â- â- â-
| 80% of 107.2 MiB

|â- â- | 90% of 107.2 MiB
|â- â- | 100% of 1
07.2 MiB
Chromium Headless Shell 143.0.7499.4 (playwright build v1200) downloaded to C:\Users
\pablo\AppData\Local\ms-playwright\chromium_headless_shell-1200
Downloading Firefox 144.0.2 (playwright build v1497) from https://cdn.playwright.de
v/dbazure/download/playwright/builds/firefox/1497/firefox-win64.zip
|
| 0% of 107.1 MiB
|â- â- â- â- â- â- | 10% of 107.1 MiB
|â- â- â- â- â- â- â- | 20% of 107.1 MiB
|â- â- â- â- â- â- â- â- | 30% of 107.1 MiB
|â- â- â- â- â- â- â- â- | 40% of 107.1 MiB
|â- â- â- â- â- â- â- â- | 50%
of 107.1 MiB
|â- â- â- â- â- â- â- | 60% of 107.1 MiB
|â- â- â- â- â- â- â- | 70% of 107.1 MiB
|â- â- â- â- â- â- â- | 80% of 107.1 MiB
|â- â- â- â- â- â- â- | 90% of 107.1 MiB
|â- â- â- â- â- â- â- | 100% of 1
07.1 MiB
Firefox 144.0.2 (playwright build v1497) downloaded to C:\Users\pablo\AppData\Local
\ms-playwright\firefox-1497
Downloading Webkit 26.0 (playwright build v2227) from https://cdn.playwright.dev/dba
zure/download/playwright/builds/webkit/2227/webkit-win64.zip
|
| 0% of 58.2 MiB
|â- â- â- â- â- â- | 10% of 58.2 MiB
|â- â- â- â- â- â- â- | 20% of 58.2 MiB
|â- â- â- â- â- â- â- | 30% of 58.2 MiB
|â- â- â- â- â- â- â- | 40% of 58.2 MiB
|â- â- â- â- â- â- â- | 50%

of 58.2 MiB
|â- â-
| 60% of 58.2 MiB
|â- â-
| 70% of 58.2 MiB
|â- â-
| 80% of 58.2 MiB
|â- â-
| 90% of 58.2 MiB
|â- â-
| 100% of 58.2 MiB
Webkit 26.0 (playwright build v2227) downloaded to C:\Users\pablo\AppData\Local\ms-playwright\webkit-2227
Downloading FFMPEG playwright build v1011 from https://cdn.playwright.dev/dbazure/download/playwright/builds/ffmpeg/1011/ffmpeg-win64.zip
|
1% of 1.3 MiB
|â- â- â- â- â- â-
| 11% of 1.3 MiB
|â- â- â- â- â- â- â- â- â- â-
| 21% of 1.3 MiB
|â- â- â- â- â- â- â- â- â- â-
| 30% of 1.3 MiB
|â- â- â- â- â- â- â- â- â- â-
| 40% of 1.3 MiB
|â- â- â- â- â- â- â- â- â- â-
| 50% of 1.3 MiB
|â- â- â- â- â- â- â- â- â-
| 60% of 1.3 MiB
|â- â- â- â- â- â- â- â- â-
| 70% of 1.3 MiB
|â- â- â- â- â- â- â- â- â-
| 80% of 1.3 MiB
|â- â- â- â- â- â- â- â- â-
| 90% of 1.3 MiB
|â- â- â- â- â- â- â- â- â-
| 100% of 1.3 MiB
FFMPEG playwright build v1011 downloaded to C:\Users\pablo\AppData\Local\ms-playwright\ffmpeg-1011
Downloading Winldd playwright build v1007 from https://cdn.playwright.dev/dbazure/download/playwright/builds/winldd/1007/winldd-win64.zip
|â- â- â- â- â-
| 10% of 0.1 MiB

```
|â- â-  
| 22% of 0.1 MiB  
|â- â-  
| 35% of 0.1 MiB  
|â- â-  
|â- â- | 48% of 0.1 MiB  
|â- â-  
| 54%  
of 0.1 MiB  
|â- â-  
| 67% of 0.1 MiB  
|â- â-  
| 79% of 0.1 MiB  
|â- â-  
| 85% of 0.1 MiB  
|â- â-  
| 98% of 0.1 MiB  
|â- â-  
| 100% of  
0.1 MiB  
Winldd playwright build v1007 downloaded to C:\Users\pablo\AppData\Local\ms-playwright\winldd-1007  
(node:16856) [DEP0169] DeprecationWarning: `url.parse()` behavior is not standardized and prone to errors that have security implications. Use the WHATWG URL API instead. CVEs are not issued for `url.parse()` vulnerabilities.  
(Use `node --trace-deprecation ...` to show where the warning was created)  
(node:14980) [DEP0169] DeprecationWarning: `url.parse()` behavior is not standardized and prone to errors that have security implications. Use the WHATWG URL API instead. CVEs are not issued for `url.parse()` vulnerabilities.  
(Use `node --trace-deprecation ...` to show where the warning was created)  
(node:14940) [DEP0169] DeprecationWarning: `url.parse()` behavior is not standardized and prone to errors that have security implications. Use the WHATWG URL API instead. CVEs are not issued for `url.parse()` vulnerabilities.  
(Use `node --trace-deprecation ...` to show where the warning was created)  
(node:16332) [DEP0169] DeprecationWarning: `url.parse()` behavior is not standardized and prone to errors that have security implications. Use the WHATWG URL API instead. CVEs are not issued for `url.parse()` vulnerabilities.  
(Use `node --trace-deprecation ...` to show where the warning was created)  
(node:21748) [DEP0169] DeprecationWarning: `url.parse()` behavior is not standardized and prone to errors that have security implications. Use the WHATWG URL API instead. CVEs are not issued for `url.parse()` vulnerabilities.  
(Use `node --trace-deprecation ...` to show where the warning was created)  
(node:2140) [DEP0169] DeprecationWarning: `url.parse()` behavior is not standardized and prone to errors that have security implications. Use the WHATWG URL API instead. CVEs are not issued for `url.parse()` vulnerabilities.  
(Use `node --trace-deprecation ...` to show where the warning was created)
```

```
In [2]: df = pd.read_csv('No_show_data.csv')  
df
```

Out[2]:

	PatientId	AppointmentID	Gender	ScheduledDay	AppointmentDay	Age	Name
0	2.987250e+13	5642903	F	2016-04-29T18:38:08Z	2016-04-29T00:00:00Z	62	
1	5.589980e+14	5642503	M	2016-04-29T16:08:27Z	2016-04-29T00:00:00Z	56	
2	4.262960e+12	5642549	F	2016-04-29T16:19:04Z	2016-04-29T00:00:00Z	62	M,
3	8.679510e+11	5642828	F	2016-04-29T17:29:31Z	2016-04-29T00:00:00Z	8	
4	8.841190e+12	5642494	F	2016-04-29T16:07:23Z	2016-04-29T00:00:00Z	56	
...
110522	2.572130e+12	5651768	F	2016-05-03T09:15:35Z	2016-06-07T00:00:00Z	56	
110523	3.596270e+12	5650093	F	2016-05-03T07:27:33Z	2016-06-07T00:00:00Z	51	
110524	1.557660e+13	5630692	F	2016-04-27T16:03:52Z	2016-06-07T00:00:00Z	21	
110525	9.213490e+13	5630323	F	2016-04-27T15:09:23Z	2016-06-07T00:00:00Z	38	
110526	3.775120e+14	5629448	F	2016-04-27T13:30:56Z	2016-06-07T00:00:00Z	54	

110527 rows × 14 columns

In [3]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 110527 entries, 0 to 110526
Data columns (total 14 columns):
 #   Column            Non-Null Count  Dtype  
--- 
 0   PatientId        110527 non-null   float64
 1   AppointmentID    110527 non-null   int64  
 2   Gender            110527 non-null   object  
 3   ScheduledDay      110527 non-null   object  
 4   AppointmentDay    110527 non-null   object  
 5   Age               110527 non-null   int64  
 6   Neighbourhood    110527 non-null   object  
 7   Scholarship       110527 non-null   int64  
 8   Hipertension      110527 non-null   int64  
 9   Diabetes          110527 non-null   int64  
 10  Alcoholism        110527 non-null   int64  
 11  Handcap           110527 non-null   int64  
 12  SMS_received      110527 non-null   int64  
 13  No-show           110527 non-null   object  
dtypes: float64(1), int64(8), object(5)
memory usage: 11.8+ MB
```

```
In [4]: print(df['No-show'].value_counts())
```

```
No-show
No     88208
Yes    22319
Name: count, dtype: int64
```

```
In [ ]:
```

```
#standardize columns
df.columns = [c.strip().lower().replace('-', '_') for c in df.columns]

#Target: no_show = 1/showed up = 0
df['no_show_flag'] = df['no_show'].map({'Yes': 1, 'No': 0})

#Drop NA

df = df.dropna(subset = ['appointmentday', 'scheduledday'])

df.no_show_flag.value_counts(normalize = True)
```

```
Out[5]: no_show_flag
0    0.798067
1    0.201933
Name: proportion, dtype: float64
```

```
In [ ]:
```

```
In [6]: df['appointmentday'] = pd.to_datetime(df['appointmentday']).dt.tz_localize(None)
df['scheduledday'] = pd.to_datetime(df['scheduledday']).dt.tz_localize(None)
```

```

# Days between scheduling and appointment
df['days_wait'] = (df['appointmentday'] - df['scheduledday']).dt.days

# Appointment weekday (Monday = 0)
df['appt_weekday'] = df['appointmentday'].dt.weekday

# Appointment Hour
df['appt_hour'] = df['appointmentday'].dt.hour

# Age Buckets
df['age_bucket'] = pd.cut(
    df['age'],
    bins=[-1, 12, 25, 40, 60, 120],
    labels=['child', 'youth', 'young_adult', 'adult', 'senior']
)

drop_cols = ['patientid', 'appointmentday', 'scheduledday', 'no_show']
df_model = df.drop(columns=[c for c in drop_cols if c in df.columns])

df_model

```

Out[6]:

	appointmentid	gender	age	neighbourhood	scholarship	hipertension	diabetes
0	5642903	F	62	JARDIM DA PENHA	0	1	0
1	5642503	M	56	JARDIM DA PENHA	0	0	0
2	5642549	F	62	MATA DA PRAIA	0	0	0
3	5642828	F	8	PONTAL DE CAMBURI	0	0	0
4	5642494	F	56	JARDIM DA PENHA	0	1	1
...
110522	5651768	F	56	MARIA ORTIZ	0	0	0
110523	5650093	F	51	MARIA ORTIZ	0	0	0
110524	5630692	F	21	MARIA ORTIZ	0	0	0
110525	5630323	F	38	MARIA ORTIZ	0	0	0
110526	5629448	F	54	MARIA ORTIZ	0	0	0

110527 rows × 15 columns

In []:

```
In [7]: # No_show rate by days_wait bucket
df_model['days_wait_bucket'] = pd.cut(df_model['days_wait'], bins=[-1, 0, 3, 7, 14, 365], labels=[-1, 0, 3, 7, 14])
no_show_by_wait = df_model.groupby('days_wait_bucket')['no_show_flag'].mean()

print(no_show_by_wait)

no_show_by_wait.plot(kind='bar')
plt.ylabel('No show rate')
plt.title('No show rate by wait time')
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()

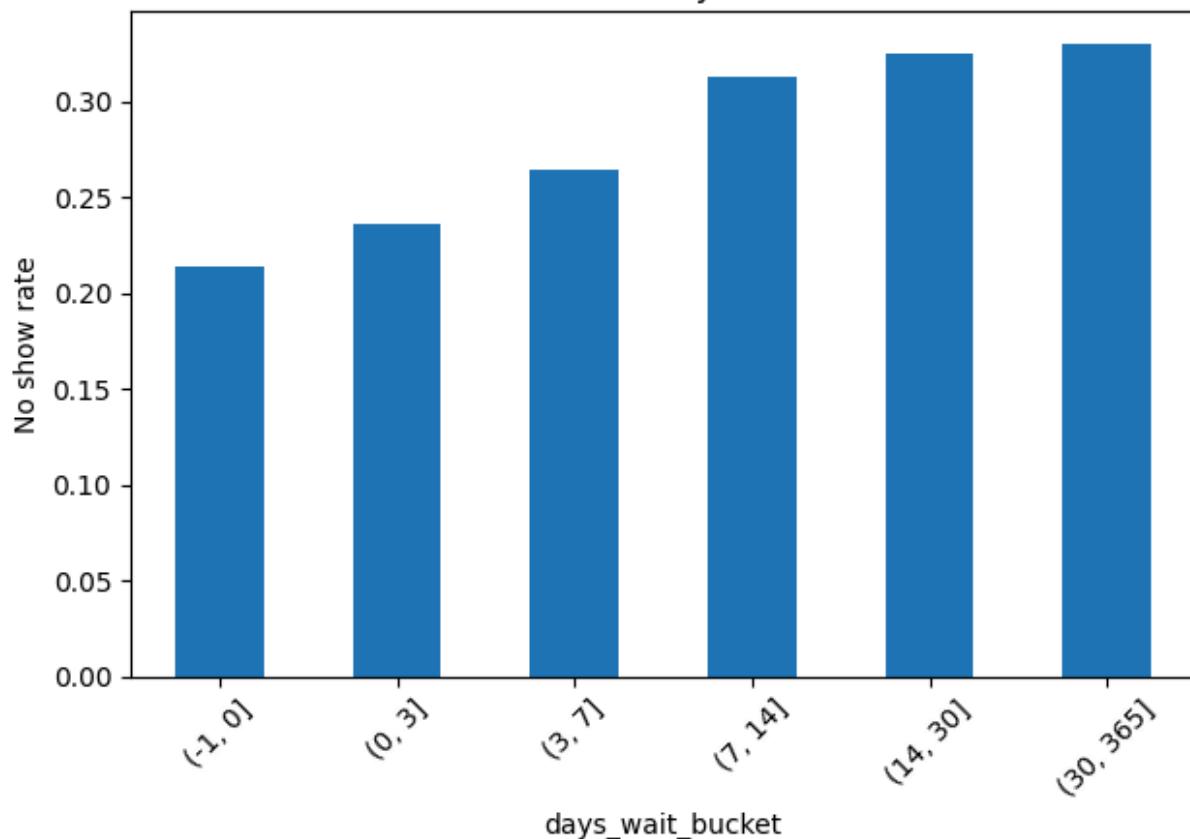
# No show by weekday
no_show_by_weekday = df_model.groupby('appt_weekday')['no_show_flag'].mean()

no_show_by_weekday.plot(kind='bar')
plt.ylabel('No show rate')
plt.title('No show rate by weekday')
plt.show()
```

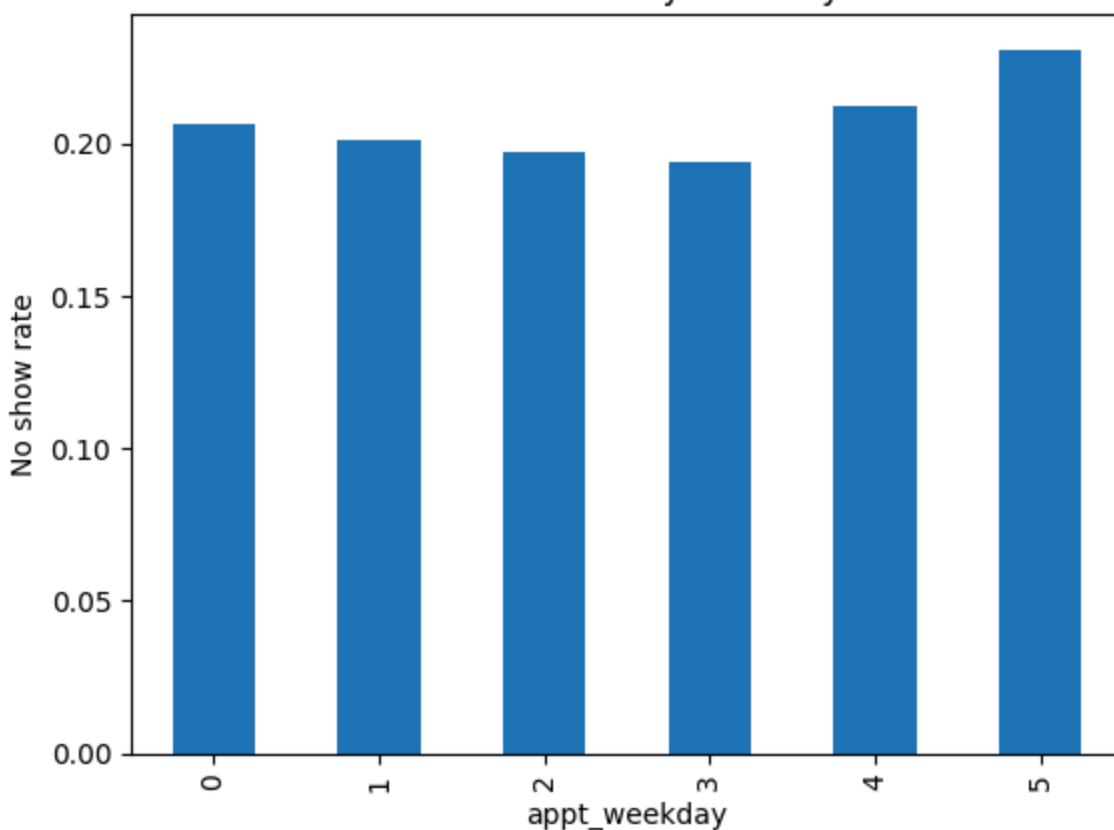
C:\Users\pablo\AppData\Local\Temp\ipykernel_20840\663844668.py:4: FutureWarning: The default of observed=False is deprecated and will be changed to True in a future version of pandas. Pass observed=False to retain current behavior or observed=True to adopt the future default and silence this warning.

```
    no_show_by_wait = df_model.groupby('days_wait_bucket')['no_show_flag'].mean()
days_wait_bucket
(-1, 0]      0.213505
(0, 3]       0.235697
(3, 7]       0.264706
(7, 14]      0.312254
(14, 30]     0.325212
(30, 365]    0.330208
Name: no_show_flag, dtype: float64
```

No show rate by wait time



No show rate by weekday



In []:

```
In [8]: ## Train/Test Split

target = 'no_show_flag'

numeric_feature = ['age', 'days_wait', 'appt_weekday', 'appt_hour', 'scholarship',
numeric_feature = [f for f in numeric_feature if f in df_model.columns]

categorical_features = ['gender', 'neighbourhood', 'age_bucket']
categorical_features = [f for f in categorical_features if f in df_model.columns]

x = df_model[numeric_feature + categorical_features]
y = df_model[target]

x_train, x_test, y_train, y_test = train_test_split(x, y, test_size = 0.2, random_s
```

```
In [ ]:
```

```
In [16]: from sklearn.preprocessing import OneHotEncoder

#Numeric pipeline

numeric_transformer = Pipeline(steps = [
    ('imputer', SimpleImputer(strategy = 'most_frequent')),
    ('encoder', OneHotEncoder(handle_unknown = 'ignore'))
])

categorical_transformer = Pipeline(steps = [
    ('imputer', SimpleImputer(strategy = 'most_frequent')),
    ('encoder', OneHotEncoder(handle_unknown = 'ignore'))
])

preprocessor = ColumnTransformer(
    transformers = [
        ('num', numeric_transformer, numeric_feature),
        ('cat', categorical_transformer, categorical_features)
    ]
)
```

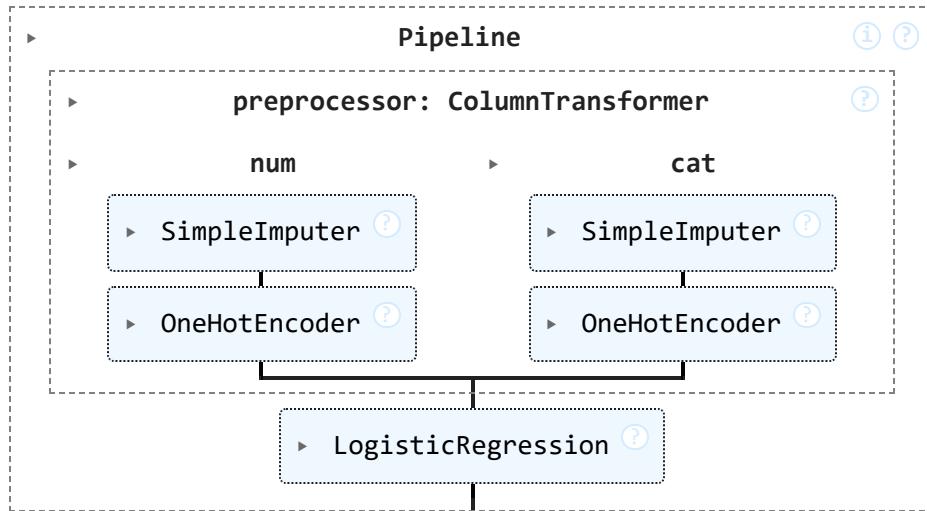
```
In [ ]:
```

```
In [17]: #Logistic Regression

log_reg_clf = Pipeline(steps = [
    ('preprocessor', preprocessor),
    ('model', LogisticRegression(max_iter = 200, n_jobs = -1))
])

log_reg_clf.fit(x_train, y_train)
```

Out[17]:



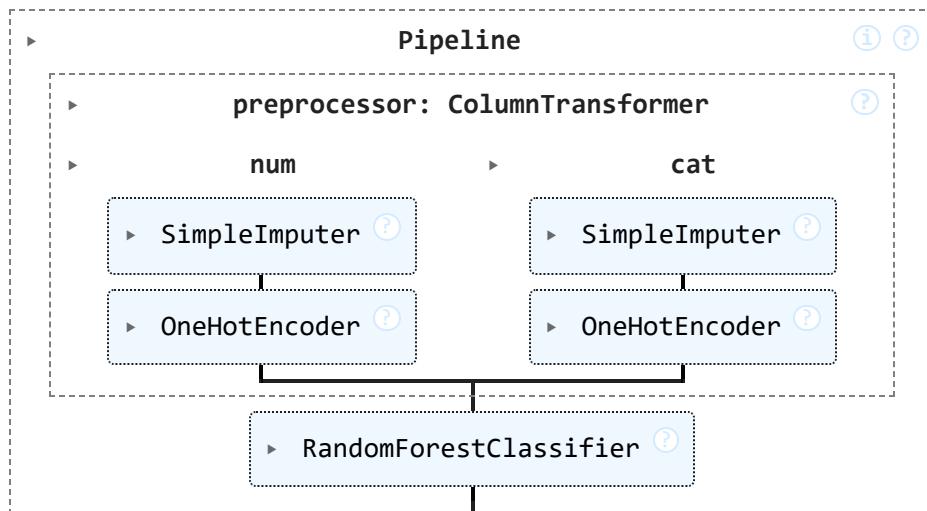
In []:

In [18]: #Random Forest

```
rf_clf = Pipeline(steps = [
    ('preprocessor', preprocessor),
    ('model', RandomForestClassifier(
        n_estimators = 300,
        max_depth = None,
        random_state = 42,
        n_jobs = -1))])

rf_clf.fit(x_train, y_train)
```

Out[18]:



In []:

In [33]: #xgboost

```
xgb_clf = Pipeline(steps = [
    ('preprocessor', preprocessor),
    ('model', XGBClassifier(
        n_estimators = 400,
```

```

        max_depth = 5,
        subsample = 0.8,
        colsample_bytree = 0.8,
        random_state = 42,
        n_jobs = -1,
        eval_metric = 'logloss',
        use_label_encoder = False))
])

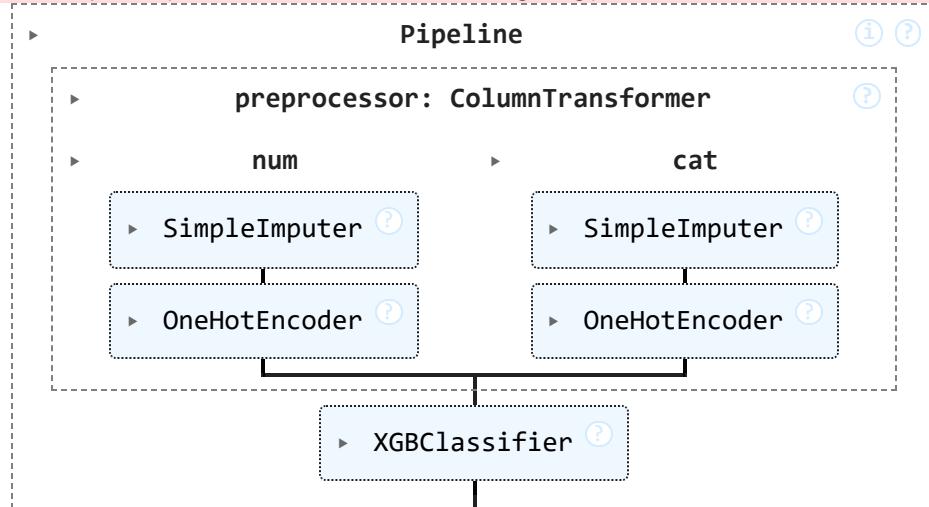
xgb_clf.fit(x_train,y_train)

```

C:\Users\pablo\AppData\Local\ABF Certificate\python\python\Lib\site-packages\xgboost\training.py:199: UserWarning: [21:16:13] WARNING: C:\actions-runner_work\xgboost\xgboost\src\learner.cc:790:
Parameters: { "use_label_encoder" } are not used.

```
bst.update(dtrain, iteration=i, fobj=obj)
```

Out[33]:



In []:

In [34]: #Evaluation Function

```

def evaluate_model(name, model, x_test, y_test):
    y_pred = model.predict(x_test)
    y_prob = model.predict_proba(x_test)[:,1]

    acc = accuracy_score(y_test, y_pred)
    prec = precision_score(y_test, y_pred)
    rec = recall_score(y_test, y_pred)
    f1 = f1_score(y_test, y_pred)
    auc = roc_auc_score(y_test, y_prob)

    print(f'== {name} ==')
    print(f'Accuracy:{acc:.3f}')
    print(f'Precision:{prec:.3f}')
    print(f'Recall:{rec:.3f}')
    print(f'F1 Score: {f1:.3f}')
    print(f'ROC AUC: {auc:.3f}')

```

```
print()

#Confusion Matrix

cm = confusion_matrix(y_test, y_pred)
print('Confusion Matrix:')
print(cm)
print('-'*40)

#ROC CURVE
fpr, tpr, thresholds = roc_curve(y_test, y_prob)
plt.plot(fpr, tpr, label = name)
plt.plot ([0,1], [0,1], linestyle = '--', color = 'black')
plt.xlabel('false positive rate')
plt.ylabel('true positive rate')
plt.title('ROC Curve')

plt.legend()
```

In []:

```
In [35]: plt(figsize=(7,5)

evaluate_model('Logistic Regression', log_reg_clf, x_test, y_test)
evaluate_model('Random Forest', rf_clf, x_test, y_test)
evaluate_model('XGBoost', xgb_clf, x_test, y_test)

plt.show()
```

==== Logistic Regression====

Accuracy:0.799

Precision:0.500

Recall:0.019

F1 Score: 0.036

ROC AUC: 0.727

Confusion Matrix:

```
[[17586    83]
 [ 4354    83]]
```

==== Random Forest====

Accuracy:0.794

Precision:0.461

Recall:0.160

F1 Score: 0.237

ROC AUC: 0.730

Confusion Matrix:

```
[[16839    830]
 [ 3728    709]]
```

==== XGBoost====

Accuracy:0.796

Precision:0.460

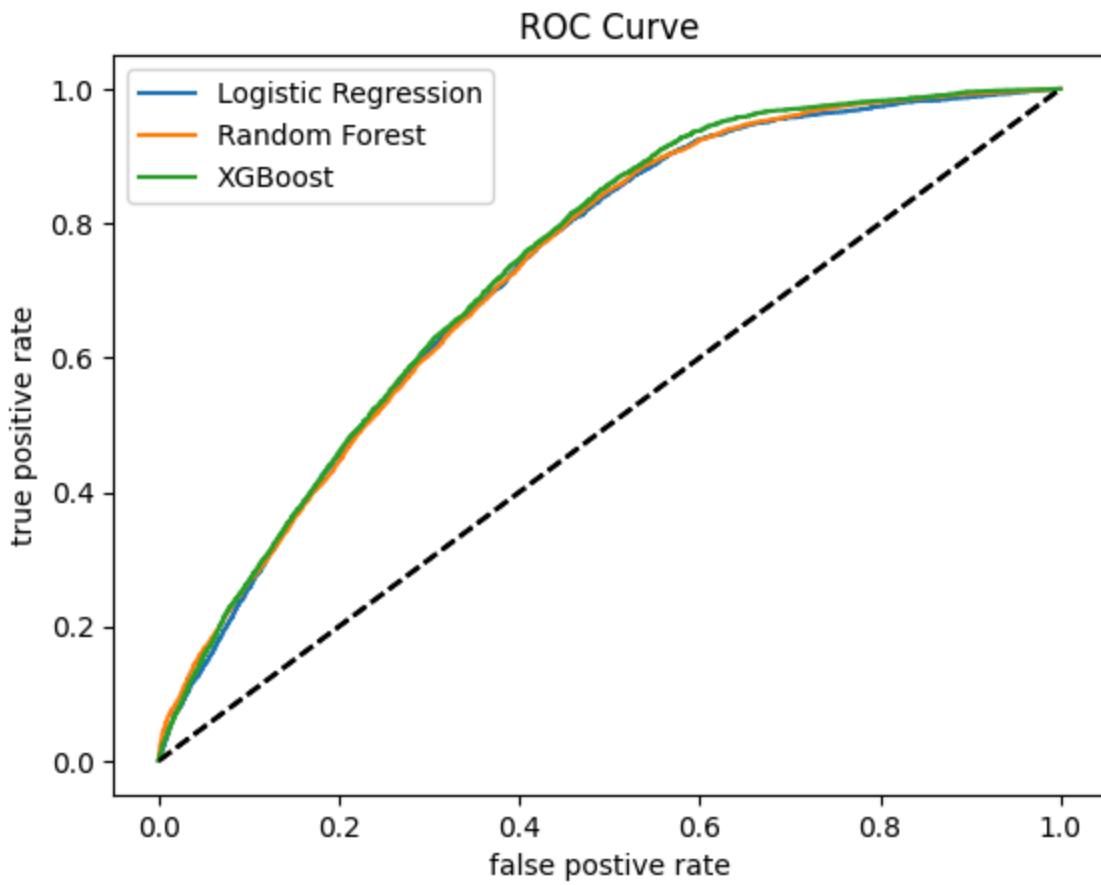
Recall:0.094

F1 Score: 0.155

ROC AUC: 0.736

Confusion Matrix:

```
[[17182    487]
 [ 4022    415]]
```



In []:

```
# All feature names in the correct order
all_feature_names = preproc.get_feature_names_out()

# Importances from RF
importances = rf_model.feature_importances_

fi_rf = (
    pd.DataFrame({
        'feature': all_feature_names,
        'importance': importances
    })
    .sort_values('importance', ascending=False) # NOTE: bool, not string
)

print(fi_rf.head(20))

# Plot top 20
fi_rf.head(20).plot(kind='barh', x='feature', y='importance')
plt.gca().invert_yaxis()
plt.title('Top 20 Feature Importances RF')
plt.tight_layout() # typo fixed
plt.show()
```

	feature	importance
105	num_days_wait_-1	0.060926
253	cat_gender_F	0.019242
254	cat_gender_M	0.019180
234	num_appt_weekday_1	0.018985
235	num_appt_weekday_2	0.018591
236	num_appt_weekday_3	0.016168
237	num_appt_weekday_4	0.016019
233	num_appt_weekday_0	0.015904
293	cat_neighbourhood_JARDIM CAMBURI	0.012247
298	cat_neighbourhood_MARIA ORTIZ	0.011192
251	num_sms_received_0	0.010467
252	num_sms_received_1	0.010461
314	cat_neighbourhood_RESISTÊNCIA	0.009832
265	cat_neighbourhood_CENTRO	0.008829
112	num_days_wait_6	0.007981
291	cat_neighbourhood_ITARARÉ	0.007863
333	cat_neighbourhood_TABUAZEIRO	0.007770
294	cat_neighbourhood_JARDIM DA PENHA	0.007470
263	cat_neighbourhood_BONFIM	0.007417
264	cat_neighbourhood_CARATOÍRA	0.007213



In []:

In [63]: #Prediction Function

```
def predict_no_show_prob (model, sample_dict):
    sample_df = pd.DataFrame ([sample_dict])
    prob = model.predict_proba(sample_df)[0,1]
```

```
return prob

patient_example = {
    'age' : 15,
    'days_wait': 7,
    'appt_weekday': 0,
    'appt_hour' : 14,
    'scholarship': 0,
    'hypertension': 0,
    'diabetes': 0,
    'alcoholism': 0,
    'handcap':0,
    'sms_received':0,
    'gender': 'M',
    'neighbourhood': (),
    'age_bucket': 'young_adult'
}

prob = predict_no_show_prob(rf_clf, patient_example)

print(f'Predicited no show probability: {prob:.2%}')
```

Predicited no show probability: 36.22%

In []:

In []: